

AI-DLC

The AI-DLC Programme

How the assessment workshop became a full service — structure, modules, and the decisions behind them.

[Programme Overview · Draft for review](#)

What this is

AI-DLC is a D55 service that helps mid-market engineering organisations actually get value from AI-assisted development — not a generic AI maturity assessment, but specifically **how to use AI across the development lifecycle**: tooling, specs, process, testing, governance, people, and ROI.

This document captures a recent step up: the free assessment workshop we built is **not the offering** — it is the first step of a larger programme. What follows is the programme structure and the reasoning behind each decision, so it can be read quickly and reacted to.

The core reframe

We were selling the diagnostic. We should be selling the **transformation**. A free one-hour diagnostic is a lead magnet, not a service. The value to a client is not “you found out where you stand” — it is “you got to AI-native engineering, you can prove it with numbers, and you own the capability to keep going without us.”

The two-minute pitch

“We assess where your engineering org really is with AI — free. We teach only the gaps that matter. Then we embed a hybrid team alongside your people, ship real work, and prove the ROI with hard numbers. Then we leave — and you keep the capability and the operating manual to run it yourselves.”

The programme: four phases

The assessment decides which phases and modules a given client actually needs. Not everyone does all of it — but where the gaps warrant it, the destination is **Phase 2 (embed)**, because that is where value is actually created.



Throughout, the Client Operating Manual is written section by section — the asset the client keeps.

What each phase does

PHASE	WHAT HAPPENS	PRODUCES
0 — Assess	The 1-hour workshop we've built. 8 dimensions, current vs target, radar + gap analysis. Free.	The starting picture, and the decision on which modules/phases to recommend.
1 — Teach	Targeted workshops, chosen by the gaps. A prescription, not a fixed menu. Clustered into four modules.	The “how you close each gap” know-how, plus the middle sections of the manual.
2 — Prove	The 8-Week Prove-It made real: a hybrid team (D55 + their people) ships live work, metrics baselined from day one.	Shipped work, a proven metrics delta, and capability transferred into the team. The centre of gravity.
3 — Scale	Expand the proven pattern across teams. Hand over the completed manual. Reduce D55 presence.	An org running AI-native engineering on its own. The exit is a feature, not a bug.

Two runbooks (a deliberate distinction)

“Runbook” was doing two jobs. Splitting it removed real confusion.

INTERNAL

D55 Delivery Playbook

Our repeatable sequence for running an engagement: assess → teach → embed → roll out. D55 IP. How a consultant knows what to do next. **Never handed over.**

THE KEPT ASSET

Client Operating Manual

The client's own playbook for AI-native engineering — their standards, tooling, guardrails, metrics, ways of working. **The artefact they own and run from after we leave.**

How they relate

Our Delivery Playbook **drives** the engagement; the engagement **produces and populates** the Client Operating Manual. Each phase we deliver writes another section of the client's manual. By the end it is complete and it is theirs.

The four workshop modules

Eight assessment dimensions would make eight workshops — a bloated menu. Clustering them into four modules keeps the service tight. Each maps to a chunk of the radar and writes one section of the Client Operating Manual.

MODULE	COVERS (DIMENSIONS)	AUDIENCE	WRITES
1. Leadership & the Investment Case	Leadership & Mandate · Metrics & ROI	Execs, sponsor, CFO/PE	§1 Mandate & Measurement
2. The AI-First Delivery Lifecycle	Tooling · Specification & Design · Delivery Process	Engineering, POs/BAs	§2 How We Build
3. Shipping Safely	Testing & Quality · Governance, Security & Compliance	Eng leads, security, compliance	§3 Shipping Safely
4. People & Change	Team Adaptation & Skills (Four Fears, middle-layer, hybrid team model)	Leadership, POs/BAs, HR	§4 People & Operating Model

Why the embed is a phase, not a module

The hybrid embedded team is the destination the modules point toward — a team operating over a multi-month engagement, not a time-boxed workshop. Demoting it to a module would shrink the highest-value part of the service into a half-day slot. It stays Phase 2. *How* to run such a team is taught in Module 4; *doing* it is Phase 2.

The hybrid embed, and the handover

Phase 2 is **not** “D55 comes and does the work.” It is a deliberately hybrid team — D55 engineers plus the client's own people on one squad — because capability only transfers if their people are on the team, learning by doing. The handover is planned from day one: D55 presence tapers as the client takes over.

The **8-Week Prove-It** (Assess wk 1–2, Prove wk 3–6, Translate wk 7–8) is the core proof cycle inside the embed — enough to baseline, ship, and demonstrate a metrics delta. A full embed typically runs longer as it tapers. Quote it as “8-week proof, then taper.”

- › **D55 + their people → handover** — hybrid squad, client staff upskill in-flight, D55 exits leaving a self-sufficient team (preferred).

- **D55-only → handover** — where the client can't spare people up front; weaker capability transfer, so we steer toward the hybrid shape.

How we drive maximum value for the client

The client doesn't want workshops or a runbook. They want three outcomes: ship faster, prove it with numbers, and don't make a mess doing it. Five principles organise the service around those.

1. Value is created in the doing, not the assessing

Assessment and workshops are setup. The value lands when a hybrid team is shipping real work with metrics baselined from day one. Get to shipped work fast and instrument it heavily.

2. Transfer capability, don't create dependency

The cleanest differentiator from a body-shop. We embed *and deliberately teach*, leaving their people able to run it. Capability transfer is an explicit, measured deliverable — not a side effect.

3. De-risking is value, especially to the buyer

The CTO's real fear is a mess at 5x speed — insecure code, compliance exposure, EU AI Act. “You can now go fast *safely*” is what lets leadership authorise scaling at all.

4. Solve the problem they don't know they have

The middle-layer bottleneck is high-value because it's non-obvious. Surfacing and fixing it makes us look like we've been somewhere they haven't.

5. Leave a flywheel

The end state is self-sustaining: metrics prove ROI → justifies investment → more adoption → more metrics. The manual institutionalises it so it survives after we leave.

Commercial shape (for discussion)

Moving from à la carte to a programme with a natural path. The existing component pricing still stands — this packages it.

TIER	INCLUDES	ROUGH SHAPE
Assess	Phase 0 only	Free (the lead magnet)
Assess + Teach	Phase 0 + chosen modules	Per module (~£5–15k, priced by duration)
Full programme	Phase 0 → Teach → Embed → Scale	Programme price; embed dominates (~£15–25k/engineer/month)

Open decisions for Rhys

A handful of calls are needed before this goes to a prospect. The structure above is designed to make them concrete rather than open-ended.

- › **Does the free assessment stay free**, or convert to paid discovery for serious prospects? (Changes the sales motion. Recommendation: start free to build trust and pipeline; add paid discovery once we have proof points.)
- › **Programme pricing** — one price, or a committed path with stage gates the client can stop at?
- › **External name** — “AI-DLC” works internally but not customer-facing (candidates already drafted).
- › **Manual format** — generated branded document (reuses existing tooling) vs a living tool the client maintains.

Where this sits

The programme structure, four modules, module schema, and Client Operating Manual outline are drafted and have passed a cold two-perspective critique (CTO 4/5, Marketing 3/5 — PASS). Remaining polish (an embed scoping one-pager, starter templates per module) is noted and authorable. The next real build is tooling: a skill that consumes the modules plus a client's assessment scores to generate a tailored write-up and populate their operating manual.