

ESG DELIVERY — CONFIDENTIAL BRIEFING

Current State & Path to Faster Delivery

Prepared for Jonathan ahead of the meeting with Lynsey

Confidential — D55 internal briefing

In one minute

- › D55 delivered four major workstreams — Financial Management, Instalment Plans, Debt Config, Debt Management — on time, to tight deadlines, running our own process.
- › The next tranche (16 initiatives) is now run under a new model: D55 developers embedded into an ESG squad, ESG leads owning design and decisions, our design-first process removed.
- › Under that model, pace has dropped. Current gut-feel estimates point to ~end-September against an end-July target — for work smaller than Debt Management, which we delivered in ~2 months with 9 developers.
- › The variable that changed is the process, not the people. The fastest route back to pace is to reinstate design-first, split the work by technical area (not business area), assign owners, then AI-accelerate from it.
- › Opportunity: ESG want to learn how we break work down for AI. We're in the process of productising exactly that — route it into a value-adding engagement rather than giving it away piecemeal.
- › Bigger picture: ESG absorbing our ways of working is the goal of our engagement model, not a threat. Treat the meeting as genuine fact-finding to uncover where else D55 can add value.

1. Current State

D55 has delivered four major workstreams to date, each on time and under tight deadlines, running our own squads and design-first process.

Delivered (D55-led, design-first process):

WORKSTREAM	OUTCOME
Financial Management	Delivered on time
Instalment Plans	Delivered on time
Debt Config	Delivered on time
Debt Management	~2 months, 9 developers — larger scope than the current tranche

Delivery quality held up under scrutiny: on the Debt/IP defect review, 127 defects, 97% closed, and ~60% self-raised by D55 before ESG testing — a team reviewing and testing its own work.

- › In flight: 16 initiatives, currently at business-requirements level only — overlapping, not yet scoped into technical design.
- › Refinement only just completed, with half of July already gone.

- › Estimate basis has changed to gut-feel, with no technical design underpinning it.

What changed operationally

- › D55 developers embedded into an ESG squad; ESG leads own the big-ticket initiatives and increasingly the overall process.
- › D55's design-first process removed — no upfront technical design, no story/subtask decomposition, no Monte Carlo forecast.
- › Autonomy removed — design and delivery decisions now sit with ESG leads.
- › Each developer owns their own initiative design → weak coordination across overlapping initiatives and integration rework late in the cycle.

2. Internal Team


The team running the current 16 initiatives is 5 developers: 3 from D55 and 2 from ESG.

- › 2 ESG developers own the main big-ticket initiatives and are taking on more of the process.
- › 3 D55 developers actively embedded, with reduced scope to make meaningful design impact.
- › 3 further D55 developers waiting on workstreams that start end of next week.
- › 2 further D55 developers working on less involved areas of the codebase.

The point to land (calmly)

The current setup under-uses D55's strengths. Our value was never interchangeable coding capacity — it was the delivery system (design-first plus AI acceleration) that produced the earlier results. Embedded as individual contributors under another process — without a view of the wider delivery, against a constantly moving baseline, and with no technical-area ownership — that advantage is lost. Not through any lack of ability, but through the absence of oversight to guide everyone in the same direction.

3. Deadlines: Expected vs Realistic



Bar positions are proportional to calendar time (July–September). The gap between the end-July target and late-September reality is the story.

- Refinement complete — mid-July. This is where we are now.
- ESG target — end of July.
- Development — ~178 dev-days across 5 developers $\approx$ 7 dev-weeks; forecast dev-complete early September, and only if perfectly parallel with zero rework.
- Testing, defect fixes and PR merges to main follow on — realistic completion late September.

How the realistic estimate builds up

- ~178 dev-days estimated across the 16 initiatives (gut-feel, no design), and excluding system testing.
- Team of 5 $\rightarrow$ ~35–36 dev-days $\approx$ ~7 dev-weeks of pure build, if perfectly parallel with zero rework.
- Excludes testing, defect-fixing, and merging feature-branch PRs into main — a non-trivial tail on every prior workstream.
- With refinement only just finished mid-July and no design to parallelise cleanly, the honest landing point is end of September.

The contrast that makes the point

Debt Management was larger, delivered in ~2 months with 9 developers under our design-first process. This tranche is smaller, yet forecast at ~2+ months with 5 developers (effectively 3 D55) under the new process.

Same-or-worse duration for less scope. The only variable that changed is the process.

A note on estimate confidence

Previously we forecast dates from technical design → user stories → subtasks → Monte Carlo simulation. Those forecasts were defensible. The current numbers are gut-feel without a design, so their confidence interval is wide. We cannot stand behind a precise date under an estimation approach that isn't ours.

4. Key Challenges

All of these trace back to one root cause: the switch away from design-first.

- › No upfront technical design — removes the artefact that makes AI acceleration possible and reliable forecasting achievable.
- › Sequential refinement dependency — work waits on individual sessions rather than being decomposed once, upfront.
- › Gut-feel estimates and ad-hoc ambiguity resolution mid-flight, rather than designed out early.
- › Fragmented design ownership — per-developer initiative design drives communication overhead and integration rework.
- › No technical-area owners. Debt Management had clear owners for front-end, back-end logic & validation, API endpoints & handlers, and the debt pipeline — each with clear accountability, clean communication lines, and very little rework, because cross-cutting design decisions were made with a view of the whole. The current per-initiative (business-area) split has none of that.
- › AI works better scoped to a specific technical area than to a business area spanning many. Technical-area ownership is what lets AI deliver faster, with better testing and less rework.
- › No comparable metrics — without our decomposition we can't produce the like-for-like throughput data that would evidence the pace drop.
- › Autonomy removed — decision latency when every meaningful call routes through ESG leads.

The AI angle that matters most (D55's own published insight)

When developers go 5x with AI, the bottleneck moves from building to specifying — "no design specs at 5x speed = failure." The current model has increased AI pressure on delivery while removing the design step that makes AI pay off, so the constraint has moved to the middle layer (design, coordination, refinement) — exactly where the slowdown is showing.

5. Our Plan to Bring It In Sooner

The fastest path is not "more people" or vaguely "more AI." It is restoring the conditions that produced our earlier pace.



The sequence matters: design first is the enabler for everything downstream.

- › Reinstate design-first: a technical design across the 16 initiatives up front, resolving the overlaps once. This enables everything else.
- › Split by technical area and assign owners: front-end, back-end logic & validation, API endpoints & handlers, and pipeline — re-establishing accountability, clear communication lines, and overall control, as we had on Debt Management.
- › AI-accelerate from the design: with clean design and defined contracts, large sections can be one-shot with Claude/Copilot — where the throughput multiplier actually lives.
- › Phase the delivery: carving into phases (already on ESG's agenda) plays to decomposition — sequence by dependency, ship value earlier.
- › Align accountability with control: happy to be measured on delivery where we own the levers; under shared ownership, timescale accountability is genuinely shared.
- › Right-size the team post-design: adding developers before the work is decomposed adds overhead and rework, not speed. Design first, then scale to the shape of the work.

Responding to the "reduce timescales" options

ESG's agenda raises three levers. Our honest position on each:

OPTION	D55 POSITION
More resource	Won't help linearly on undesigned, overlapping work — risks worsening integration rework. Design first, assign technical-area owners, then add people to a decomposed backlog organised by technical area, not business area.

OPTION	D55 POSITION
Use of AI	The real lever — but it requires the upfront design. This is the honest case for reinstating design-first.
Phasing	Supportive — it aligns with how we decompose.

6. Where Else D55 Can Help

We don't currently have visibility of Lynsey's own priorities, so this is genuine fact-finding. The areas below are hypotheses to test in the conversation, not claims to present.

The reframe to hold onto

ESG wanting to absorb our ways of working is the intended outcome of our engagement model — embed, prove, then enable, not a permanent dependency. That demand is a market for our highest-value offering, not a threat to defend against. Approach the whole conversation as discovering where that capability transfer is most valuable to her.

Value areas to explore

Grounded in what we know of ESG's domain and the signals to date. Items marked (inference) are our read, not documented fact — test them, don't assert them.

AREA	WHAT WE'D EXPLORE
Specification & decomposition at AI pace	The middle-layer bottleneck — Lynsey's own domain. When devs go 5x, the constraint moves to specifying. Help ESG's product/BA function spec and decompose at pace.
Reviewer / lead capacity	"Damian is spread too thin." Codified standards and PR-review enablement relieve senior-reviewer load so it stops being a bottleneck.
Predictable forecasting	Design → stories → Monte Carlo forecasting vs gut-feel. Predictable dates often matter more to a product leader than fast ones. (inference)
Domain-depth delivery	Pricing and Billing, Debt, Instalment Plans, Comms Hub — hard-won context D55 already holds, backed by excellent breadth and depth of energy-sector knowledge. Continued delivery of high-value initiatives here. (inference)
Greenfield done right	Alex's own point: the Debt Service greenfield was a missed clean-slate chance. An opening for D55 to lead clean service design under agreed standards. (inference)

Questions to uncover it

Each question maps to a D55 engagement model once we hear the answer (workshops, forecasting/DSD, embedded team, enablement).

- What outcomes are you measured on this year — roadmap, predictability, cost, quality?
- Where does delivery most often stall for you — specifying, reviewing, testing, integrating?
What's the constraint you can't just hire your way out of?

- Which piece of upcoming work worries you most?
- Six months from now, what would make you say bringing D55 in was worth it?

7. How to Play the Conversation with Lynsey

Strategic framing — Jonathan's call, offered as counsel.

- › Listen first. Lynsey is product/outcome-focused and more senior than the day-to-day delivery politics. Understand her pain points and goals before positioning.
- › Partner, not defence. Avoid leading with "the timescales aren't our responsibility" — true in fairness terms, but it invites "then we'll take it fully in-house." Frame as "align accountability with control" and pivot to value.
- › Anchor on the contrast, not the complaint: "smaller scope, same duration, only the process changed" is fair, factual, and lands without blaming individuals.
- › Move up the value stack. As embedded individuals we are commoditised and easy to offboard; as the people who bring the 5x method we are not.
- › Turn "teach us your AI method" into an opportunity. We're in the process of productising this (the AI-DLC programme). Offer it as a structured enablement engagement — protects the method, secures future work, strengthens the relationship, answers the AI question.
- › Be measured on the "how." Share the principle (design-first unlocks AI) freely; keep the detailed reusable playbook as the commercial asset it is.
- › Jonathan's real goal: reinforce the relationship and secure future work. Every point should serve that, not win a delivery-process argument.

Open questions to confirm before the meeting

- › Is end-of-July a hard ESG commitment to a customer, or an internal target?
- › What is Lynsey's own top priority — a date, a cost, a capability transfer, or de-risking?
- › How explicit do we want to be about the handover signals — name it, or let Lynsey lead there?
- › Appetite to formally propose the AI-DLC enablement engagement, or plant the seed only?