

MODULE 2

A standard AI tooling configuration (with shared skills / MCP connections identified)

The AI-First Delivery Lifecycle

Workshop deliverable · template

Purpose

The minimum viable, shared AI-assisted development setup every engineer gets — so the team moves from “licences provisioned” to “embedded in daily flow.” Standardise it, check it into version control, and give it to new joiners on day one.

The standard config

LAYER	WHAT GOOD LOOKS LIKE	YOUR CHOICE
IDE + AI assistant	Agreed editor + assistant (Copilot / Claude / Cursor / Kiro)	___
Shared skills / rules	Team prompt library + coding-standard rules, versioned in the repo	___
MCP connections	Assistant wired to Jira, GitHub, and key data sources	___
Context files	Repo-level guidance (architecture, conventions) the assistant reads	___
PR-review settings	AI first-pass review configured; human bar defined for AI code	___
Secrets / DLP	No secrets in prompts; approved endpoints only (see Module 3)	___

The test of “embedded”

A new joiner clones the repo and is productive with the assistant the same day — no one-off setup, no “ask Dave how he configured his.”

Adoption checklist

- > Config lives in the repo, not on individual machines.
- > There's a named owner who keeps it current.
- > Onboarding doc points new joiners at it.
- > Usage is visible (who's actually using it, and how).